


Mean over runs · shaded ± 1 std across runs · smoothed 15 gens · converged = last 50 gens



converged mean ± 1 std across runs

Evolution (EA only)
2.894 $\pm$ 0.105 (n = 100)

Learning (EA + in-life RL)
2.89 $\pm$ 0.103 (n = 100)

Pure RL (no EA)
2.714 $\pm$ 0.192 (n = 100)