




converged mean ± 1 std across runs

— Evolution (EA only)
0.7123 $\pm$ 0.297 (n = 100)

— Learning (EA + in-life RL)
0.7919 $\pm$ 0.184 (n = 100)

— Pure RL (no EA)
0.579 $\pm$ 0.234 (n = 100)